

# HW5-Coding(5)

January 11, 2024

## 1 Homework 5: Convolutional neural network (30 points)

In this part, you need to implement and train a convolutional neural network on the CIFAR-10 dataset with PyTorch. ### What is PyTorch?

PyTorch is a system for executing dynamic computational graphs over Tensor objects that behave similarly as numpy ndarray. It comes with a powerful automatic differentiation engine that removes the need for manual back-propagation.

### 1.0.1 Why?

- Our code will now run on GPUs! Much faster training. When using a framework like PyTorch or TensorFlow you can harness the power of the GPU for your own custom neural network architectures without having to write CUDA code directly (which is beyond the scope of this class).
- We want you to be ready to use one of these frameworks for your project so you can experiment more efficiently than if you were writing every feature you want to use by hand.
- We want you to stand on the shoulders of giants! TensorFlow and PyTorch are both excellent frameworks that will make your lives a lot easier, and now that you understand their guts, you are free to use them :)
- We want you to be exposed to the sort of deep learning code you might run into in academia or industry.

• GPU PyTorch TensorFlow GPU CUDA

•

• TensorFlow PyTorch )

• ### How can I learn PyTorch?

Justin Johnson has made an excellent [tutorial](#) for PyTorch.

You can also find the detailed [API doc](#) here. If you have other questions that are not addressed by the API docs, the [PyTorch forum](#) is a much better place to ask than StackOverflow.

Install PyTorch and Skorch.

```
[2]: !pip install -q torch skorch torchvision torchtext
```

```
[2]: import torch
import torch.nn as nn
import torch.nn.functional as F
```

```
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
import skorch
import sklearn
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

## 1.1 0. Tensor Operations (5 points)

Tensor operations are important in deep learning models. In this part, you are required to get familiar to some common tensor operations in PyTorch.

### 1.1.1 1) Tensor squeezing, unsqueezing and viewing

Tensor squeezing, unsqueezing and viewing are important methods to change the dimension of a Tensor, and the corresponding functions are [torch.squeeze](#), [torch.unsqueeze](#) and [torch.Tensor.view](#). Please read the documents of the functions, and finish the following practice.

```
[9]: # x is a tensor with size being (3, 2)
x = torch.Tensor([[1, 2],
                  [3, 4],
                  [5, 6]])

print(x.shape)

# Add two new dimensions to x by using the function torch.unsqueeze
x = torch.unsqueeze(torch.unsqueeze(x, -1), 1)
print(x.shape)

# Remove the two dimensions just added by using the function torch.squeeze
x = torch.squeeze(torch.squeeze(x, -1), 1)
print(x.shape)

# x is now a two-dimensional tensor, or in other words a matrix. Now use the
→function torch.Tensor.view and change x to a one-dimensional vector with
→size being (6).
x = x.view(-1)
print(x.shape)
```

```
torch.Size([3, 2])
torch.Size([3, 1, 2, 1])
torch.Size([3, 2])
torch.Size([6])
```

### 1.1.2 2) Tensor concatenation and stack

Tensor concatenation and stack are operations to combine small tensors into big tensors. The corresponding functions are [torch.cat](#) and [torch.stack](#). Please read the documents of the functions,

and finish the following practice.

```
[10]: # x is a tensor with size being (3, 2)
x = torch.Tensor([[1, 2], [3, 4], [5, 6]])

# y is a tensor with size being (3, 2)
y = torch.Tensor([[-1, -2], [-3, -4], [-5, -6]])

# Our goal is to generate a tensor z with size as (2, 3, 2), and z[0,:,:] = x,
↪ z[1,:,:] = y.

# Use torch.stack to generate such a z
# pass
z = torch.stack([x, y])
print(z[0,:,:])
# Use torch.cat and torch.unsqueeze to generate such a z
# pass
z = torch.cat([x.unsqueeze(0), y.unsqueeze(0)], dim = 0 )
print(z[1,:,:])

tensor([[1., 2.],
        [3., 4.],
        [5., 6.]])
tensor([[-1., -2.],
        [-3., -4.],
        [-5., -6.]])
```

### 1.1.3 3) Tensor expansion

Tensor expansion is to expand a tensor into a larger tensor along singleton dimensions. The corresponding functions are `torch.Tensor.expand` and `torch.Tensor.expand_as`. Please read the documents of the functions, and finish the following practice.

```
[11]: # x is a tensor with size being (3)
x = torch.Tensor([1, 2, 3])

# Our goal is to generate a tensor z with size (2, 3), so that z[0,:,:] = x,
↪ z[1,:,:] = x.

# [TO DO]
# Change the size of x into (1, 3) by using torch.unsqueeze.
# pass
x = torch.unsqueeze(x, 0)
print(x.shape)

# [TO DO]
# Then expand the new tensor to the target tensor by using torch.Tensor.expand.
# pass
```

```
z = x.expand(2, -1)
print(z.shape)
```

```
torch.Size([1, 3])
torch.Size([2, 3])
```

#### 1.1.4 4) Tensor reduction in a given dimension

In deep learning, we often need to compute the mean/sum/max/min value in a given dimension of a tensor. Please read the document of [torch.mean](#), [torch.sum](#), [torch.max](#), [torch.min](#), [torch.topk](#), and finish the following practice.

```
[12]: # x is a random tensor with size being (10, 50)
x = torch.randn(10, 50)

# Compute the mean value for each row of x.
# You need to generate a tensor x_mean of size (10), and x_mean[k, :] is the
# ↪ mean value of the k-th row of x.
# pass
x_mean = x.mean(dim=1)
print(x_mean[3, ])

# Compute the sum value for each row of x.
# You need to generate a tensor x_sum of size (10).
# pass
x_sum = x.sum(dim=1)
print(x_sum.shape)

# Compute the max value for each row of x.
# You need to generate a tensor x_max of size (10).
# pass
x_max, _ = x.max(dim=1)
print(x_max.shape)

# Compute the min value for each row of x.
# You need to generate a tensor x_min of size (10).
# pass
x_min, _ = x.min(dim=1)
print(x_min.shape)

# Compute the top-5 values for each row of x.
# You need to generate a tensor x_mean of size (10, 5), and x_top[k, :] is the
# ↪ top-5 values of each row in x.
# pass
x_xtop, _ = torch.topk(x, k=5, dim=1, largest=True, sorted=True)
print((x_xtop.shape))
```

```
tensor(0.1405)
torch.Size([10])
torch.Size([10])
torch.Size([10])
torch.Size([10, 5])
```

## 1.2 Convolutional Neural Networks

Implement a convolutional neural network for image classification on CIFAR-10 dataset.

CIFAR-10 is an image dataset of 10 categories. Each image has a size of 32x32 pixels. The following code will download the dataset, and split it into `train` and `test`. For this question, we use the default validation split generated by Skorch.

```
[3]: train = torchvision.datasets.CIFAR10("./data", train=True, download=True)
test = torchvision.datasets.CIFAR10("./data", train=False, download=True)
```

Files already downloaded and verified

Files already downloaded and verified

The following code visualizes some samples in the dataset. You may use it to debug your model if necessary.

```
[6]: def plot(data, labels=None, num_sample=5):
    n = min(len(data), num_sample)
    for i in range(n):
        plt.subplot(1, n, i+1)
        plt.imshow(data[i], cmap="gray")
        plt.xticks([])
        plt.yticks([])
        if labels is not None:
            plt.title(labels[i])

train.labels = [train.classes[target] for target in train.targets]
plot(train.data, train.labels)
```



### 1.2.1 1) Basic CNN implementation

Consider a basic CNN model

- It has 3 convolutional layers, followed by a linear layer.
- Each convolutional layer has a kernel size of 3, a padding of 1.
- ReLU activation is applied on every hidden layer.

Please implement this model in the following section. The hyperparameters is then be tuned and you need to fill the results in the table. CNN

- 3
- 3 1
- ReLU

**a) Implement convolutional layers (10 Points)** Implement the initialization function and the forward function of the CNN.

```
[3]: class CNN(nn.Module):
    def __init__(self, channels):
        super(CNN, self).__init__()
        # implement parameter definitions here
        # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
        # pass
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=channels,
→kernel_size=3, padding=1)
        self.conv2 = nn.Conv2d(in_channels=channels, out_channels=channels, 3,
→padding=1)
        self.conv3 = nn.Conv2d(in_channels=channels, out_channels=channels, 3,
→padding=1)
        self.fc = nn.Linear(channels * 32 * 32, 10)
        # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
    def forward(self, images):
        # implement the forward function here
        # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
        # pass
        images = images.float()
        images = F.relu(self.conv1(images))
        images = F.relu(self.conv2(images))
        images = F.relu(self.conv3(images))
        images = images.view(images.size(0), -1)
        images = self.fc(images)
        # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
        return images
```

**b) Tune hyperparameters** Train the CNN model on CIFAR-10 dataset. We can tune the number of channels, optimizer, learning rate and the number of epochs for best validation accuracy.

[4]:

```

# implement hyperparameters, you can select and modify the hyperparameters by
↳ yourself here.

optimize = [torch.optim.SGD, torch.optim.Adam]
learning_rate = [1e-3]
channel = [16, 32, 64]

train_data_normalized = torch.Tensor(train.data/255)
train_data_normalized = train_data_normalized.permute(0,3,1,2)

for l in learning_rate:
    for o in optimize:
        for c in channel:
            print(f'The channel was {c}, the learning rate was {l} and the optimizer
↳ was {str(o)}')

            cnn = CNN(channels = c)

            model = skorch.NeuralNetClassifier(cnn, criterion=torch.nn.
↳ CrossEntropyLoss,

                                                device="cuda",
                                                optimizer=o,
                                                # optimizer__momentum=0.90,
                                                lr=l,
                                                max_epochs=50,
                                                batch_size=512,
                                                callbacks=[skorch.callbacks.
↳ EarlyStopping(lower_is_better=True)])

            # implement input normalization & type cast here
            model.fit(train_data_normalized, torch.LongTensor(train.targets))

```

The channel was 16, the learning rate was 0.001 and the optimizer was <class 'torch.optim.sgd.SGD'>

epoch	train_loss	valid_acc	valid_loss	dur
1	2.3023	0.1000	2.3016	
1.8696				
2	2.3014	0.1001	2.3006	
1.4992				
3	2.3004	0.1002	2.2996	
1.6483				
4	2.2993	0.1003	2.2985	
1.4937				
5	2.2982	0.1026	2.2973	
1.4817				
6	2.2970	0.1084	2.2959	

1.6490			
7	2.2955	0.1191	2.2943
1.6139			
8	2.2939	0.1261	2.2925
1.5891			
9	2.2920	0.1352	2.2904
1.7191			
10	2.2898	0.1441	2.2879
1.6297			
11	2.2870	0.1536	2.2848
1.5851			
12	2.2836	0.1639	2.2809
1.5934			
13	2.2794	0.1746	2.2760
1.7376			
14	2.2740	0.1833	2.2698
1.5488			
15	2.2672	0.1979	2.2620
1.5256			
16	2.2587	0.2089	2.2523
1.7559			
17	2.2482	0.2197	2.2404
1.5962			
18	2.2354	0.2299	2.2260
1.6180			
19	2.2200	0.2397	2.2088
1.6428			
20	2.2017	0.2477	2.1885
1.4556			
21	2.1804	0.2568	2.1652
1.4557			
22	2.1565	0.2635	2.1396
1.4744			
23	2.1308	0.2695	2.1127
1.5492			
24	2.1045	0.2754	2.0859
1.4287			
25	2.0791	0.2785	2.0606
1.4330			
26	2.0558	0.2838	2.0380
1.5790			
27	2.0353	0.2874	2.0186
1.4708			
28	2.0179	0.2900	2.0022
1.4527			
29	2.0031	0.2937	1.9882
1.5850			
30	1.9903	0.2973	1.9762



1.5017			
31	1.9790	0.3011	1.9654
1.4780			
32	1.9687	0.3045	1.9555
1.5058			
33	1.9590	0.3079	1.9462
1.6080			
34	1.9498	0.3139	1.9373
1.4211			
35	1.9410	0.3188	1.9288
1.4571			
36	1.9324	0.3243	1.9205
1.5877			
37	1.9241	0.3280	1.9125
1.4352			
38	1.9161	0.3324	1.9048
1.4467			
39	1.9084	0.3349	1.8974
1.4845			
40	1.9010	0.3399	1.8902
1.5439			
41	1.8938	0.3413	1.8833
1.4759			
42	1.8869	0.3453	1.8766
1.4976			
43	1.8803	0.3480	1.8702
1.5401			
44	1.8740	0.3506	1.8640
1.4488			
45	1.8679	0.3526	1.8580
1.4856			
46	1.8621	0.3540	1.8523
1.6838			
47	1.8564	0.3570	1.8469
1.5008			
48	1.8510	0.3594	1.8416
1.4871			
49	1.8459	0.3609	1.8366
1.5410			
50	1.8409	0.3613	1.8318
1.5418			

The channel was 32, the learning rate was 0.001 and the optimizer was <class 'torch.optim.sgd.SGD'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	2.3021	0.1041	2.3013	
2.8145				
2	2.3004	0.1489	2.2995	

2.6793				
3	2.2986	0.1506	2.2975	
2.5704				
4	2.2965	0.1301	2.2953	2.4280
5	2.2941	0.1228	2.2925	2.5590
6	2.2911	0.1244	2.2892	3.1142
7	2.2875	0.1309	2.2850	3.2769
8	2.2829	0.1431	2.2798	2.8112
9	2.2771	0.1596	2.2732	
2.6166				
10	2.2698	0.1762	2.2647	
2.7641				
11	2.2602	0.1919	2.2536	
2.6307				
12	2.2479	0.2082	2.2394	
2.6573				
13	2.2321	0.2180	2.2213	
2.7342				
14	2.2124	0.2290	2.1989	
2.6742				
15	2.1884	0.2427	2.1725	
2.6494				
16	2.1611	0.2525	2.1431	
2.7171				
17	2.1320	0.2595	2.1131	
2.7030				
18	2.1035	0.2667	2.0848	
2.7112				
19	2.0776	0.2721	2.0598	
2.6295				
20	2.0554	0.2777	2.0386	
2.7505				
21	2.0369	0.2822	2.0212	
2.5853				
22	2.0215	0.2849	2.0067	
2.6423				
23	2.0086	0.2890	1.9944	
2.6975				
24	1.9973	0.2925	1.9836	
2.6076				
25	1.9872	0.2953	1.9738	
2.6293				
26	1.9778	0.3010	1.9647	
2.6288				
27	1.9688	0.3049	1.9561	
2.8176				
28	1.9602	0.3063	1.9477	
2.6918				

29	1.9517	0.3116	1.9395
2.6841			
30	1.9434	0.3166	1.9315
2.7937			
31	1.9351	0.3232	1.9236
2.9334			
32	1.9269	0.3274	1.9157
2.8763			
33	1.9188	0.3329	1.9079
2.8930			
34	1.9107	0.3356	1.9003
3.2919			
35	1.9029	0.3395	1.8928
3.1846			
36	1.8952	0.3441	1.8855
3.0136			
37	1.8877	0.3453	1.8785
3.0824			
38	1.8805	0.3474	1.8717
3.0156			
39	1.8736	0.3494	1.8653
2.9976			
40	1.8671	0.3527	1.8591
3.0877			
41	1.8609	0.3533	1.8533
2.9834			
42	1.8551	0.3548	1.8479
2.9826			
43	1.8496	0.3567	1.8427
2.9340			
44	1.8444	0.3586	1.8378
3.0666			
45	1.8396	0.3605	1.8333
3.0049			
46	1.8350	0.3623	1.8290
2.9624			
47	1.8308	0.3651	1.8250
2.9678			
48	1.8267	0.3667	1.8212
2.9123			
49	1.8229	0.3679	1.8177
2.8755			
50	1.8193	0.3696	1.8143
2.9986			

The channel was 64, the learning rate was 0.001 and the optimizer was <class 'torch.optim.sgd.SGD'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----

1	2.2999	0.1120	2.2971	
5.2526				
2	2.2950	0.1218	2.2919	
4.9220				
3	2.2895	0.1216	2.2857	4.8481
4	2.2826	0.1310	2.2777	
5.0554				
5	2.2736	0.1529	2.2669	
4.8972				
6	2.2613	0.1813	2.2523	
4.8617				
7	2.2446	0.2026	2.2324	
5.0537				
8	2.2222	0.2200	2.2062	
4.9247				
9	2.1936	0.2322	2.1742	
4.9192				
10	2.1604	0.2488	2.1395	
4.8985				
11	2.1265	0.2574	2.1063	
5.1471				
12	2.0949	0.2722	2.0765	
5.0504				
13	2.0660	0.2828	2.0491	
5.0806				
14	2.0390	0.2953	2.0234	
4.9728				
15	2.0135	0.3068	1.9994	
4.8894				
16	1.9901	0.3146	1.9781	
4.8898				
17	1.9698	0.3214	1.9601	
5.0013				
18	1.9527	0.3285	1.9453	
4.6879				
19	1.9386	0.3322	1.9332	
4.7185				
20	1.9267	0.3350	1.9232	
4.7069				
21	1.9166	0.3388	1.9145	
4.8730				
22	1.9077	0.3404	1.9070	
4.7222				
23	1.8998	0.3424	1.9000	
4.7212				
24	1.8925	0.3440	1.8934	
4.9435				
25	1.8857	0.3471	1.8872	

5.0120				
26	1.8794	0.3491	1.8815	
4.9071				
27	1.8734	0.3509	1.8760	
5.1758				
28	1.8677	0.3522	1.8708	
5.1380				
29	1.8623	0.3537	1.8656	
4.9481				
30	1.8571	0.3561	1.8609	
4.9800				
31	1.8521	0.3558	1.8563	5.4353
32	1.8473	0.3578	1.8518	
5.1319				
33	1.8427	0.3589	1.8472	
4.9452				
34	1.8382	0.3615	1.8425	
5.0308				
35	1.8337	0.3629	1.8381	
4.9283				
36	1.8292	0.3650	1.8337	
4.9094				
37	1.8249	0.3683	1.8296	
4.8871				
38	1.8207	0.3700	1.8255	
5.0013				
39	1.8166	0.3711	1.8214	
4.9111				
40	1.8126	0.3715	1.8175	
4.9419				
41	1.8086	0.3716	1.8135	
4.9686				
42	1.8046	0.3731	1.8096	
4.8710				
43	1.8008	0.3758	1.8058	
4.9026				
44	1.7969	0.3763	1.8021	
5.1728				
45	1.7932	0.3778	1.7985	
5.2594				
46	1.7894	0.3779	1.7948	
5.2541				
47	1.7857	0.3788	1.7912	
5.0514				
48	1.7820	0.3803	1.7876	
5.3261				
49	1.7783	0.3810	1.7841	
5.0562				

50            1.7746            0.3819            1.7804  
5.2772  
The channel was 16, the learning rate was 0.001 and the optimizer was <class  
'torch.optim.adam.Adam'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	1.8488	0.4405	1.5880	
2.1092				
2	1.5285	0.4807	1.4754	
1.4926				
3	1.4184	0.4920	1.4440	
1.4509				
4	1.3420	0.5363	1.3145	
1.5134				
5	1.2640	0.5557	1.2645	
1.6366				
6	1.2032	0.5642	1.2354	
1.4786				
7	1.1497	0.5737	1.2122	
1.4639				
8	1.1042	0.5810	1.1967	
1.6642				
9	1.0645	0.5827	1.1885	
1.4675				
10	1.0287	0.5853	1.1848	
1.5633				
11	0.9958	0.5882	1.1829	
1.6633				
12	0.9635	0.5900	1.1732	
1.4790				
13	0.9336	0.5959	1.1675	
1.5211				
14	0.9071	0.5989	1.1675	
1.5810				
15	0.8830	0.6014	1.1697	1.6498
16	0.8628	0.5991	1.1860	1.5684
17	0.8463	0.6043	1.1800	1.5647

Stopping since valid\_loss has not improved in the last 5 epochs.  
The channel was 32, the learning rate was 0.001 and the optimizer was <class  
'torch.optim.adam.Adam'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	1.8419	0.4501	1.5560	
2.7578				
2	1.4838	0.4990	1.4230	
2.7789				
3	1.3575	0.5249	1.3367	
2.9298				

4	1.2650	0.5562	1.2526	
2.7506				
5	1.1903	0.5699	1.2085	
2.7766				
6	1.1195	0.5903	1.1619	
2.7301				
7	1.0474	0.6011	1.1307	
2.8848				
8	0.9762	0.6080	1.1117	
2.7578				
9	0.9106	0.6153	1.0970	
2.7570				
10	0.8492	0.6209	1.1026	2.8909
11	0.7926	0.6262	1.1078	2.7543
12	0.7423	0.6199	1.1382	2.7523
13	0.7030	0.6220	1.1651	2.6920

Stopping since valid\_loss has not improved in the last 5 epochs.

The channel was 64, the learning rate was 0.001 and the optimizer was <class 'torch.optim.adam.Adam'>

epoch	train_loss	valid_acc	valid_loss	dur
1	1.7742	0.4654	1.4959	
4.9542				
2	1.4214	0.5256	1.3271	
5.0191				
3	1.2805	0.5337	1.3249	
5.0329				
4	1.1849	0.5756	1.2055	
4.9213				
5	1.0875	0.5998	1.1314	
4.9105				
6	1.0037	0.6161	1.1106	
5.0302				
7	0.9147	0.6194	1.1166	4.8979
8	0.8371	0.6238	1.1196	4.8899
9	0.7662	0.6238	1.1438	4.9137
10	0.7078	0.6197	1.1802	5.0204

Stopping since valid\_loss has not improved in the last 5 epochs.

Write down **validation accuracy** of your model under different hyperparameter settings. Note the validation set is automatically split by Skorch during `model.fit()`.

#channel for each layer	optimizer	SGD	Adam
16		0.3613	0.6043
32		0.3696	0.6220
64		0.3819	0.6197

### 1.2.2 2) Full CNN implementation (10 points)

Based on the CNN in the previous question, implement a full CNN model with max pooling layer.

- Add a max pooling layer after each convolutional layer.
- Each max pooling layer has a kernel size of 2 and a stride of 2.

Please implement this model in the following section. The hyperparameters is then be tuned and fill the results in the table. You are also required to complete the questions.

**a) Implement max pooling layers** Similar to the CNN implementation in previous question, implement max pooling layers.

```
[6]: class CNN_MaxPool(nn.Module):
    def __init__(self, channels):
        super(CNN_MaxPool, self).__init__()
        # implement parameter definitions here
        # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=channels,
        ↪kernel_size=3, padding=1)
        self.conv2 = nn.Conv2d(in_channels=channels, out_channels=channels, 3,
        ↪padding=1)
        self.conv3 = nn.Conv2d(in_channels=channels, out_channels=channels, 3,
        ↪padding=1)
        self.pool = nn.MaxPool2d(kernel_size=2, stride=2)
        self.fc = nn.Linear(channels * 4 * 4, 10)
        # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

    def forward(self, images):
        # implement the forward function here
        # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
        images = images.float()
        images = F.relu(self.conv1(images))
        images = self.pool(images)
        images = F.relu(self.conv2(images))
        images = self.pool(images)
        images = F.relu(self.conv3(images))
        images = self.pool(images)
        images = images.view(images.size(0), -1)
        images = self.fc(images)
        # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
        return images
```

**b) Tune hyperparameters** Based on the better optimizer found in the previous problem, we can tune the number of channels and learning rate for best validation accuracy.

[7]:



```

# implement hyperparameters, you can select and modify the hyperparameters by
↳yourself here.
learning_rate = [1e-4]
channel = [16, 32, 64]
# Select the better optimizer by the result shown in the previous problem, you
↳can select and modify it by yourself here.
better_optimizer = torch.optim.Adam

train_data_normalized = torch.Tensor(train.data/255)
train_data_normalized = train_data_normalized.permute(0,3,1,2)

for l in learning_rate:
    for c in channel:
        print(f'The channel was {c}, the learning rate was {l}')

        cnn = CNN_MaxPool(channels = c)

        model = skorch.NeuralNetClassifier(cnn, criterion=torch.nn.
↳CrossEntropyLoss,
                                         device="cuda",
                                         optimizer=better_optimizer,
                                         lr=l,
                                         max_epochs=50,
                                         batch_size=256,
                                         callbacks=[skorch.callbacks.
↳EarlyStopping(lower_is_better=True)])
        # implement input normalization & type cast here
        model.fit(train_data_normalized, torch.LongTensor(train.targets))

```

The channel was 16, the learning rate was 0.0001

epoch	train_loss	valid_acc	valid_loss	dur
1	2.2904	0.1483	2.2464	
1.6172				
2	2.1098	0.2743	2.0010	
1.1716				
3	1.9816	0.2989	1.9379	
1.1853				
4	1.9221	0.3272	1.8772	
1.3423				
5	1.8557	0.3562	1.8055	
1.2601				
6	1.7853	0.3769	1.7474	
1.2338				
7	1.7325	0.3892	1.7060	
1.3598				
8	1.6936	0.4004	1.6745	

1.2577			
9	1.6637	0.4078	1.6490
1.2484			
10	1.6397	0.4128	1.6277
1.2435			
11	1.6198	0.4195	1.6096
1.4338			
12	1.6027	0.4267	1.5941
1.2394			
13	1.5879	0.4324	1.5802
1.2567			
14	1.5748	0.4354	1.5680
1.2290			
15	1.5631	0.4398	1.5568
1.3180			
16	1.5525	0.4442	1.5468
1.2734			
17	1.5427	0.4501	1.5376
1.2090			
18	1.5336	0.4544	1.5290
1.3671			
19	1.5252	0.4571	1.5211
1.2738			
20	1.5173	0.4590	1.5136
1.1364			
21	1.5099	0.4618	1.5066
1.1105			
22	1.5030	0.4642	1.5000
1.1440			
23	1.4964	0.4670	1.4937
1.1733			
24	1.4901	0.4699	1.4879
1.1235			
25	1.4840	0.4728	1.4822
1.2580			
26	1.4782	0.4741	1.4766
1.1354			
27	1.4727	0.4750	1.4713
1.1519			
28	1.4674	0.4773	1.4663
1.1752			
29	1.4622	0.4794	1.4614
1.1096			
30	1.4571	0.4817	1.4566
1.1074			
31	1.4522	0.4844	1.4520
1.1333			
32	1.4475	0.4864	1.4476

1.3321			
33	1.4428	0.4872	1.4432
1.1371			
34	1.4382	0.4881	1.4389
1.1324			
35	1.4337	0.4901	1.4347
1.1373			
36	1.4293	0.4918	1.4305
1.1809			
37	1.4249	0.4931	1.4264
1.1323			
38	1.4207	0.4945	1.4225
1.1026			
39	1.4164	0.4954	1.4187
1.2285			
40	1.4123	0.4974	1.4148
1.1756			
41	1.4082	0.4988	1.4109
1.2780			
42	1.4042	0.5008	1.4071
1.1071			
43	1.4002	0.5028	1.4036
1.1156			
44	1.3963	0.5042	1.4000
1.0883			
45	1.3924	0.5058	1.3964
1.2562			
46	1.3886	0.5076	1.3929
1.3020			
47	1.3849	0.5098	1.3894
1.1344			
48	1.3812	0.5113	1.3860
1.1567			
49	1.3775	0.5121	1.3825
1.2061			
50	1.3740	0.5128	1.3792

1.1291

The channel was 32, the learning rate was 0.0001

epoch	train_loss	valid_acc	valid_loss	dur
1	2.2316	0.2693	2.0385	
1.4971				
2	1.9780	0.3181	1.9140	
1.4968				
3	1.8640	0.3640	1.8010	
1.6392				
4	1.7552	0.3979	1.7090	
1.5602				

5	1.6730	0.4178	1.6434	
1.5724				
6	1.6142	0.4307	1.5934	
1.5481				
7	1.5697	0.4460	1.5546	
1.4835				
8	1.5350	0.4547	1.5236	
1.4542				
9	1.5072	0.4621	1.4983	
1.5190				
10	1.4845	0.4686	1.4776	
1.5926				
11	1.4651	0.4779	1.4595	
1.4701				
12	1.4482	0.4847	1.4434	
1.5429				
13	1.4335	0.4879	1.4293	
1.4540				
14	1.4205	0.4944	1.4168	
1.4837				
15	1.4089	0.4983	1.4060	
1.4929				
16	1.3986	0.5023	1.3960	
1.5204				
17	1.3890	0.5060	1.3870	
1.5871				
18	1.3802	0.5098	1.3787	
1.4465				
19	1.3719	0.5138	1.3710	
1.5009				
20	1.3642	0.5163	1.3635	
1.4504				
21	1.3569	0.5181	1.3566	
1.4646				
22	1.3499	0.5210	1.3501	
1.5039				
23	1.3432	0.5236	1.3438	
1.4552				
24	1.3367	0.5236	1.3379	1.5805
25	1.3305	0.5253	1.3323	
1.4698				
26	1.3245	0.5282	1.3267	
1.5021				
27	1.3186	0.5308	1.3212	
1.4621				
28	1.3128	0.5326	1.3159	
1.4453				
29	1.3072	0.5350	1.3105	

1.5120				
30	1.3017	0.5373	1.3052	
1.4677				
31	1.2963	0.5387	1.3002	
1.5693				
32	1.2910	0.5421	1.2953	
1.5502				
33	1.2858	0.5435	1.2905	
1.4665				
34	1.2806	0.5449	1.2857	
1.4535				
35	1.2755	0.5483	1.2810	
1.4499				
36	1.2704	0.5496	1.2765	
1.5073				
37	1.2654	0.5503	1.2719	
1.6197				
38	1.2605	0.5509	1.2673	
1.7351				
39	1.2556	0.5519	1.2629	
1.5843				
40	1.2508	0.5542	1.2584	
1.5823				
41	1.2459	0.5563	1.2540	
1.5619				
42	1.2412	0.5571	1.2497	
1.4999				
43	1.2364	0.5590	1.2455	
1.4913				
44	1.2317	0.5615	1.2412	
1.5563				
45	1.2271	0.5629	1.2370	
1.6428				
46	1.2225	0.5647	1.2329	
1.4798				
47	1.2179	0.5668	1.2288	
1.5729				
48	1.2134	0.5682	1.2249	
1.5736				
49	1.2090	0.5688	1.2209	
1.5438				
50	1.2046	0.5702	1.2171	
1.3991				
The channel was 64, the learning rate was 0.0001				
epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	2.1464	0.3092	1.9339	
2.1533				

2	1.8240	0.3776	1.7582
2.1884			
3	1.6660	0.4231	1.6407
2.2465			
4	1.5776	0.4492	1.5584
2.2363			
5	1.5184	0.4701	1.4993
2.2044			
6	1.4741	0.4857	1.4553
2.2599			
7	1.4384	0.4956	1.4204
2.1999			
8	1.4078	0.5055	1.3921
2.1791			
9	1.3819	0.5169	1.3664
2.2640			
10	1.3591	0.5260	1.3448
2.1652			
11	1.3386	0.5312	1.3249
2.1339			
12	1.3199	0.5389	1.3074
2.2700			
13	1.3026	0.5459	1.2912
2.2839			
14	1.2866	0.5508	1.2763
2.2272			
15	1.2716	0.5563	1.2629
2.3535			
16	1.2576	0.5619	1.2501
2.3623			
17	1.2443	0.5662	1.2380
2.3617			
18	1.2315	0.5692	1.2270
2.2513			
19	1.2194	0.5740	1.2163
2.2911			
20	1.2079	0.5771	1.2062
2.2833			
21	1.1967	0.5806	1.1965
2.3257			
22	1.1860	0.5836	1.1873
2.3147			
23	1.1757	0.5866	1.1785
2.5167			
24	1.1657	0.5889	1.1701
2.2488			
25	1.1560	0.5918	1.1620
2.1878			

26	1.1466	0.5939	1.1541
2.3689			
27	1.1376	0.5969	1.1464
2.1535			
28	1.1288	0.5989	1.1391
2.1696			
29	1.1202	0.6021	1.1319
2.1288			
30	1.1118	0.6047	1.1246
2.2984			
31	1.1036	0.6073	1.1178
2.1223			
32	1.0955	0.6092	1.1111
2.1786			
33	1.0877	0.6107	1.1049
2.1347			
34	1.0802	0.6124	1.0987
2.1286			
35	1.0728	0.6135	1.0926
2.1851			
36	1.0656	0.6145	1.0871
2.1227			
37	1.0587	0.6163	1.0817
2.3853			
38	1.0518	0.6177	1.0793
2.1707			
39	1.0453	0.6226	1.0714
2.2874			
40	1.0389	0.6255	1.0667
2.2098			
41	1.0326	0.6275	1.0621
2.2192			
42	1.0266	0.6304	1.0577
2.1843			
43	1.0207	0.6320	1.0535
2.2130			
44	1.0149	0.6331	1.0496
2.4063			
45	1.0093	0.6348	1.0458
2.1557			
46	1.0037	0.6368	1.0420
2.1941			
47	0.9983	0.6388	1.0383
2.1705			
48	0.9931	0.6394	1.0347
2.2612			
49	0.9880	0.6414	1.0312
2.2446			

50            0.9829            0.6433            1.0277  
2.2944

Write down the **validation accuracy** of the model under different hyperparameter settings.

#channel for each layer	validation accuracy
16	0.5128
32	0.5702
64	0.6433

For the best model you have, test it on the test set.

```
[8]: # implement the same input normalization & type cast here
test_data_normalized = torch.Tensor(test.data/255)
test_data_normalized = test_data_normalized.permute(0,3,1,2)
test.predictions = model.predict(test_data_normalized)
sklearn.metrics.accuracy_score(test.targets, test.predictions)
```

[8]: 0.6438

How much **test accuracy** do you get? What can you conclude for the design of CNN structure and tuning of hyperparameters? (5 points)

**Your Answer:** 0.6438 It can be concluded that increasing the number of channels helps to improve the accuracy.

[ ]: